

# Scaling on AWS for the First 10 Million Users




# Scaling on AWS for the First 10 Million Users

- ME: – Solutions Architect – Amazon Web Services – [jman@amazon.com](mailto:jman@amazon.com)
- YOU: Here to learn more about scaling infrastructure on AWS
- TODAY: about best practices and things to think about when building for large scale



# So how do we scale?





WeKnowMemes

scaling on AWS



**Web**

Images

Maps

Shopping

More ▾

Search tools

About 921,000 results (0.18 seconds)

### [Auto Scaling - Amazon Web Services](#)

[aws.amazon.com/autoscaling/](https://aws.amazon.com/autoscaling/)

Auto **Scaling** allows you to automatically **scale** your Amazon **EC2** capacity up or down according to conditions you define.

You've visited this page many times. Last visit: 10/15/12

### [Auto Scaling Documentation - Amazon Web Services](#)

[aws.amazon.com/documentation/autoscaling/](https://aws.amazon.com/documentation/autoscaling/)

List of official **AWS** documentation for Auto **Scaling**, including the Developer Guide and Getting Started Guide.

You've visited this page 4 times. Last visit: 2/20/13



scaling on AWS



a lot of things to read

Web

Images

Maps

Shopping

More ▾

Search tools

About 921,000 results (0.18 seconds)

### Auto Scaling - Amazon Web Services

[aws.amazon.com/autoscaling/](https://aws.amazon.com/autoscaling/)

Auto **Scaling** allows you to automatically **scale** your Amazon **EC2** capacity up or down according to conditions you define.

You've visited this page many times. Last visit: 10/15/12

### Auto Scaling Documentation - Amazon Web Services

[aws.amazon.com/documentation/autoscaling/](https://aws.amazon.com/documentation/autoscaling/)

List of official **AWS** documentation for Auto **Scaling**, including the Developer Guide and Getting Started Guide.

You've visited this page 4 times. Last visit: 2/20/13

scaling on AWS



a lot of things to read

Web

Images

Maps

Shopping

More ▾

Search tools

About 921,000 results (0.18 seconds)

### Auto Scaling - Amazon Web Services

[aws.amazon.com/autoscaling/](https://aws.amazon.com/autoscaling/)

Auto **Scaling** allows you to automatically **scale** your Amazon **EC2** capacity up or down according to conditions you define.

You've visited this page many times. Last visit: 10/15/12

not where we want to start

### Auto Scaling Documentation - Amazon Web Services

[aws.amazon.com/documentation/autoscaling/](https://aws.amazon.com/documentation/autoscaling/)

List of official **AWS** documentation for Auto **Scaling**, including the Developer Guide and Getting Started Guide.

You've visited this page 4 times. Last visit: 2/20/13

Auto-Scaling is a tool and a destination. It's not the single thing that fixes everything.





# What do we need first?



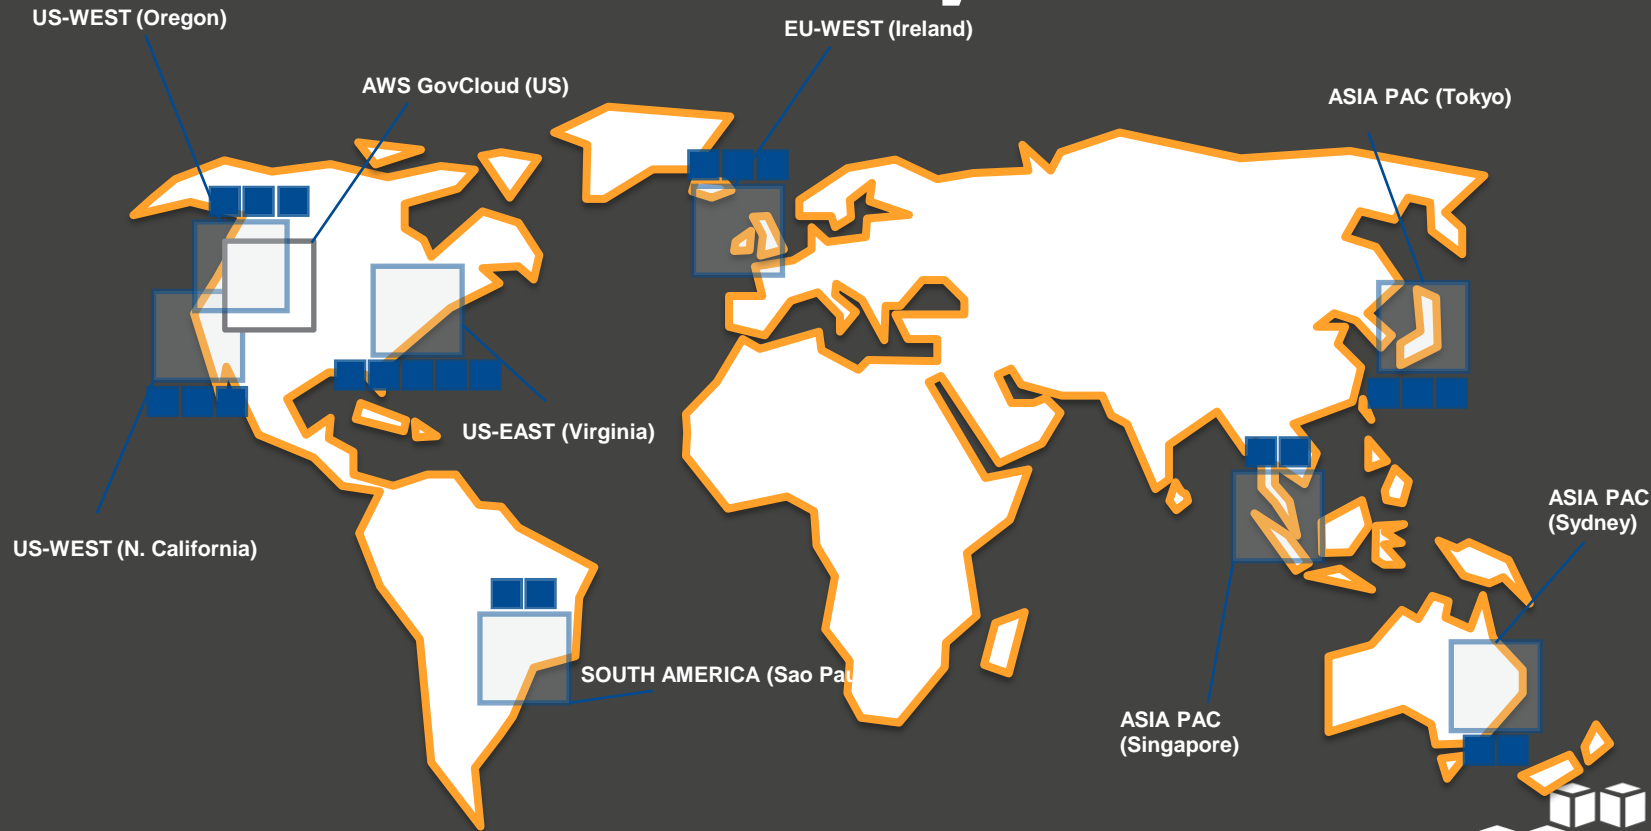
# Some basics:



# Regions



# Availability Zones



# 2003

amazon.com

- \$5.2B retail business
- 7,800 employees
- A whole lot of servers

# 2013



**Every day**, AWS adds enough  
server capacity to power that  
whole **\$5B** enterprise



Deployment & Administration

App Services

Compute

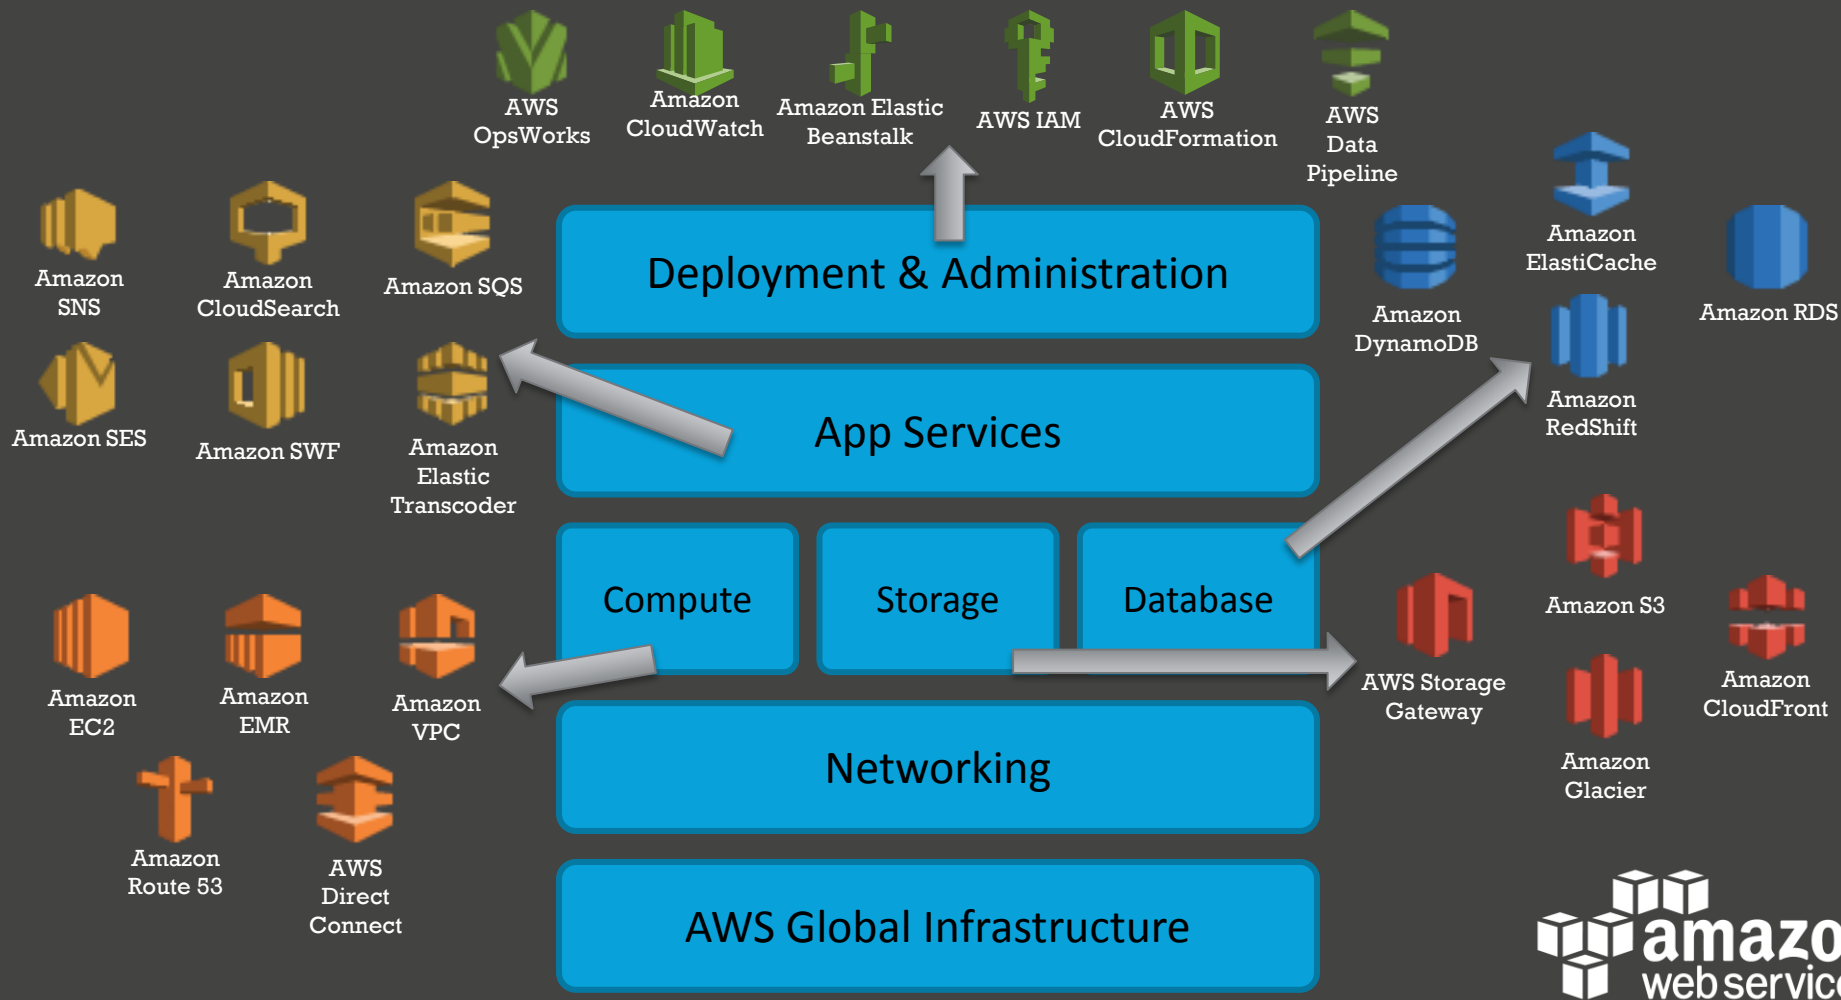
Storage

Database

Networking

AWS Global Infrastructure





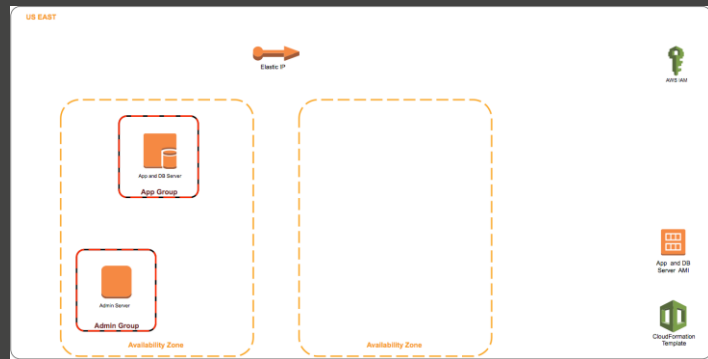
So let's start from day  
one, user one ( you ):






# Lab 1

- AWS Setup
  - security groups
  - key pairs
- Application Setup
  - Insoshi
  - MySQL




# Day One, User One:

- A single EC2 Instance
  - With full stack on this host
    - Web App
    - Database
    - Management
    - Etc.
- A single Elastic IP
- Route53 for DNS




# “We’re gonna need a bigger box”

- Simplest approach
- Can now leverage PIOPs
- High I/O instances
- High Memory instances
- High CPU instances
- High storage instances
- Easy to change instance sizes
- Will hit an endpoint eventually



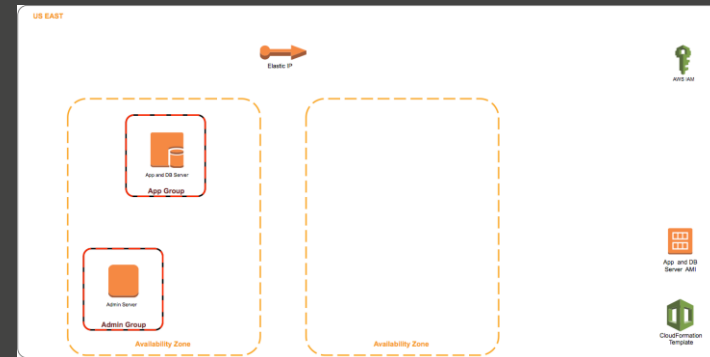
# “We’re gonna need a bigger box”

- Simplest approach
- Can now leverage PIOPs
- High I/O instances
- High memory instances
- High CPU instances
- High storage instances
- Easy to change instance sizes
- **Will hit an endpoint eventually**




# Lab 2

- Instance Resizing
  - Impact to service availability
  - Impact to data persistence




# Day One, User One:

- We could potentially get to a few hundred to a few thousand depending on application complexity and traffic
- No failover
- No redundancy
- Too many eggs in one basket



# Day One, User One:

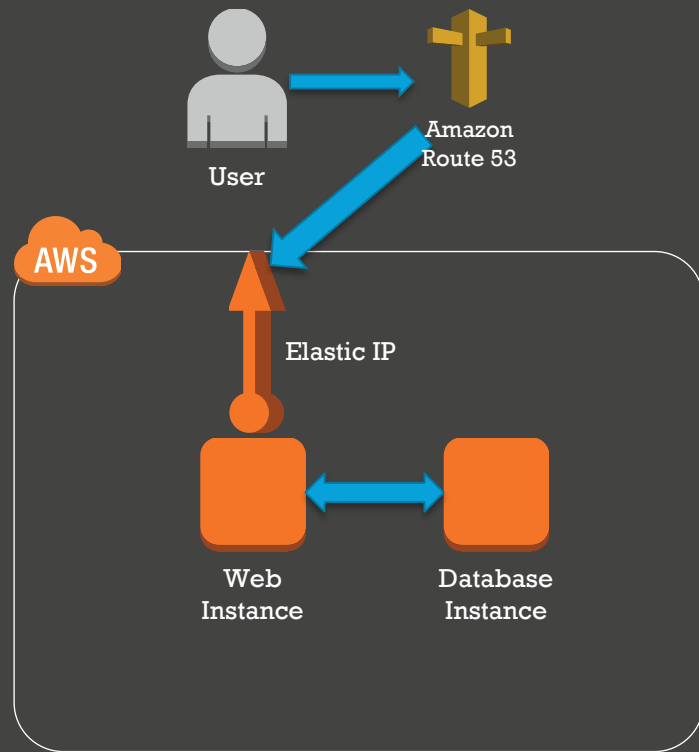
- We could potentially get to a few hundred to a few thousand depending on application complexity and traffic
- **No failover**
- **No redundancy**
- **Too many eggs in one basket**



# Day Two, User > 1:

First let's separate out our single host into more than one.

- Web
- Database
  - Make use of a database service?





# Database Options

## Self-Managed



### Database Server on Amazon EC2

Your choice of  
database running on  
Amazon EC2

Bring Your Own  
License (BYOL)

## Fully-Managed



### Amazon RDS

Microsoft SQL,  
Oracle or MySQL as  
a managed service


Flexible licensing  
BYOL or License  
Included



### Amazon DynamoDB

Managed NoSQL  
database service  
using SSD storage

Seamless  
scalability Zero  
administration



### Amazon Redshift

Massively parallel,  
petabyte-scale,  
data warehouse  
service

Fast, powerful and  
easy to scale



But how do I choose  
what DB technology I  
need? SQL? NoSQL?

Some folks won't like  
this. But...



# Start with SQL databases



But, but, but, but...

No. You don't.



# Start with SQL databases



# Why start with SQL?

- Established and well worn technology
- Lots of existing code, communities, books, background, tools, etc
- You aren't going to break SQL DBs in your first 10 million users. No really, you won't\*
- Clear patterns to scalability

\*Unless you are doing something SUPER weird with the data or MASSIVE amounts of it, even then SQL will have a place in your stack





AH HA! You said “massive  
amounts”, I will have  
massive amounts!



If your usage is such that you  
will be generating several TB  
( >5 ) of data in the first year  
OR have an incredibly data  
intensive workload you might  
need NoSQL



# Why else might you need NoSQL?

- Super low latency applications
- Metadata driven datasets
- Highly-unrelational data
- Need schema-less data constructs\*
- Massive amounts of data (again, in the TB range)
- Rapid ingest of data ( thousands of records/sec )

\*Need != “its easier to do dev without schemas”



But this is probably  
less than 90% of you



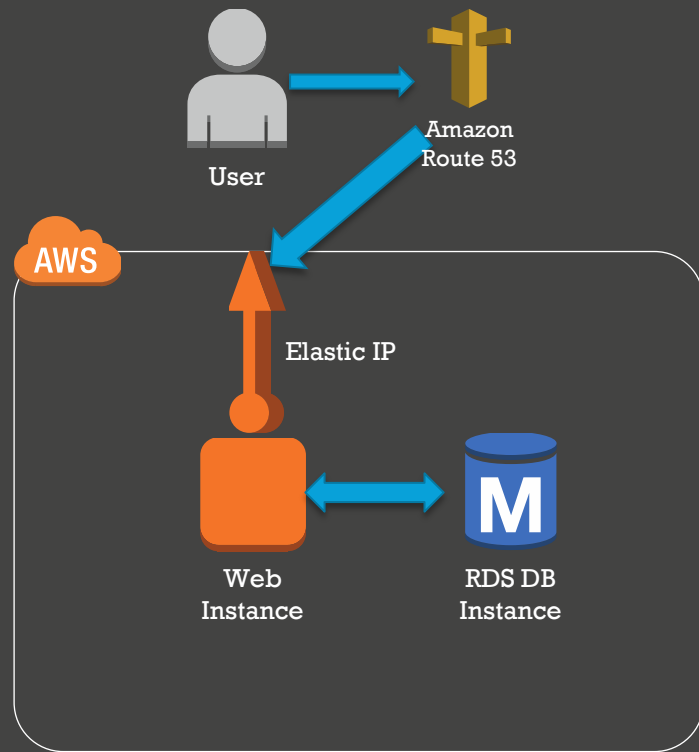
Unless everyone of you is  
building semantic/big  
data websites



# User > 100:

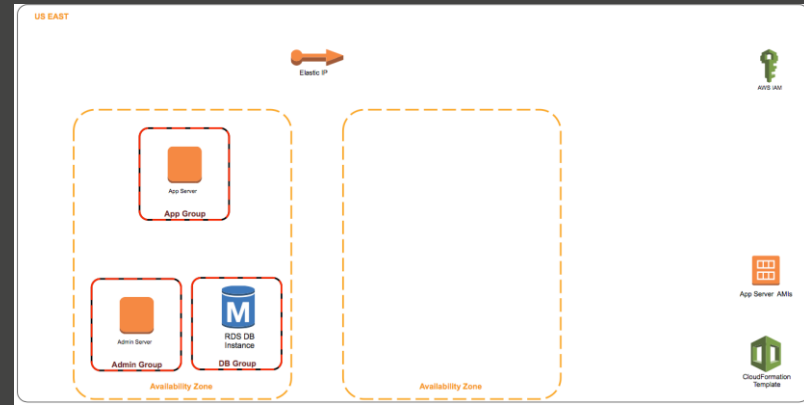
First let's separate out our single host into more than one.

- Web
- Database
  - Use RDS to make your life easier



# Lab 3 RDS

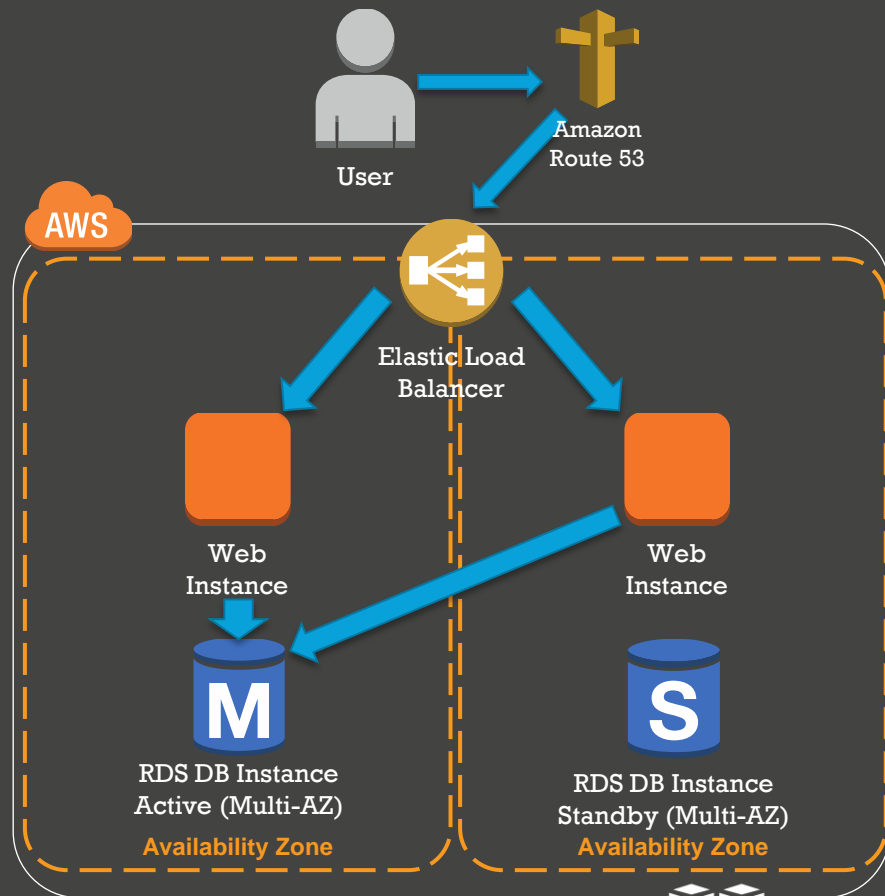
- Review
  - differences for running MySQL on RDS, if any
  - current limitation of you architecture



# User > 1000:

Next let's address our lack of failover and redundancy issues:


- Elastic Load Balancer
- Another Web Instance
  - In another Availability Zone
- Enable RDS Multi-AZ





# Elastic Load Balancing

- Create highly scalable applications
- Distribute load across EC2 instances in multiple availability zones

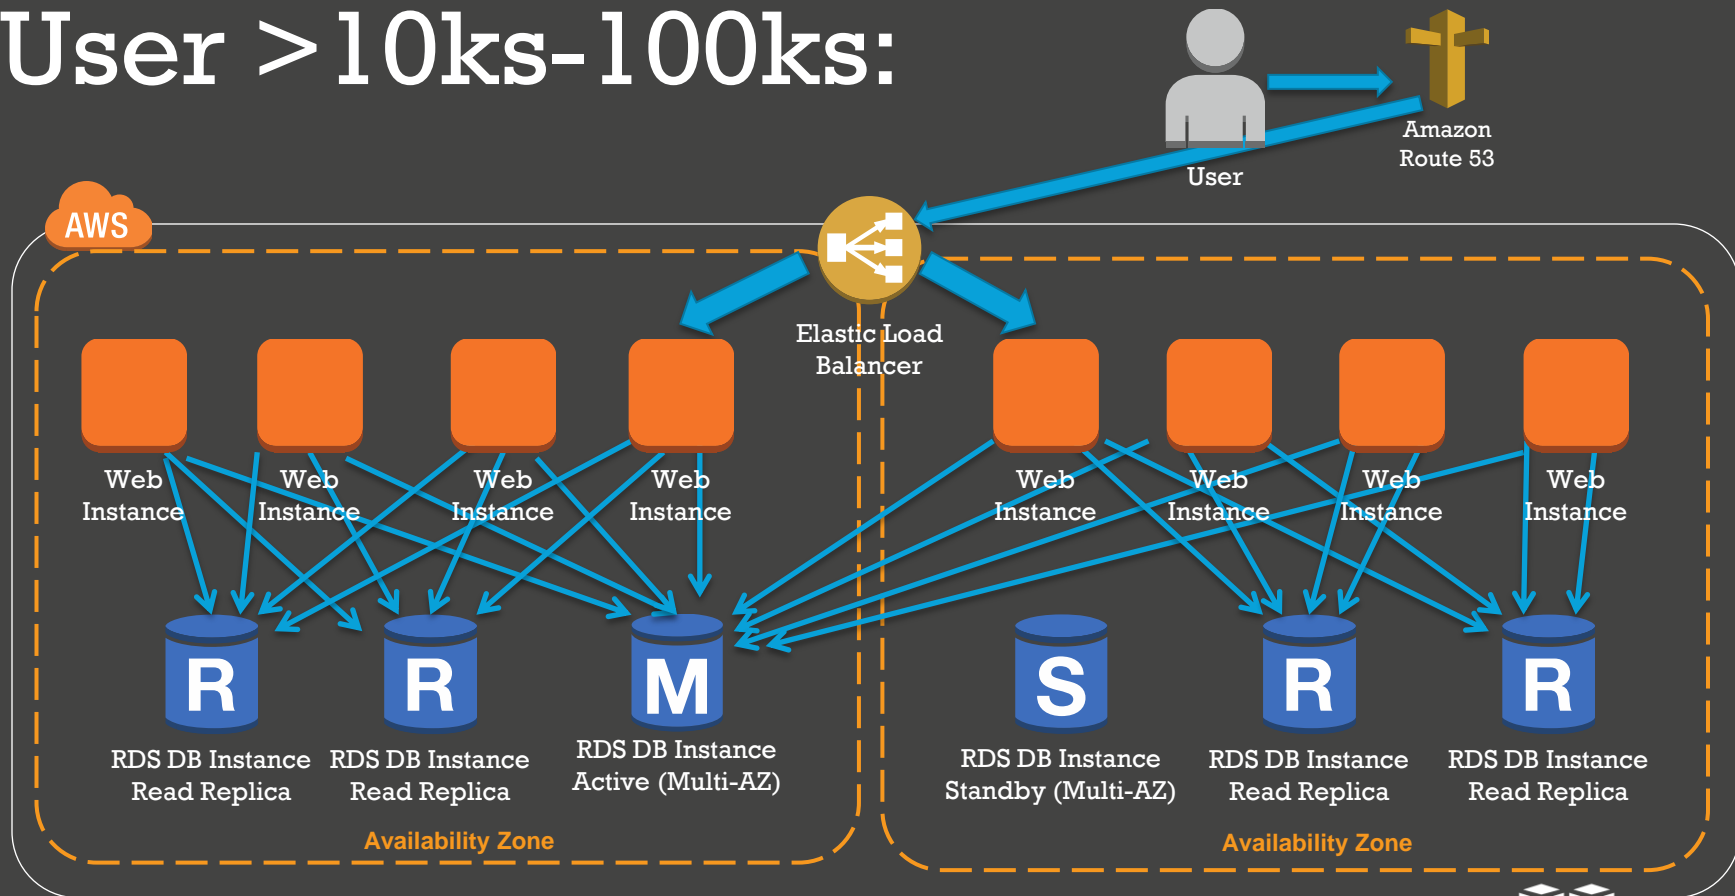


Feature	Details
<b>Available</b>	Load balance across instances in multiple Availability Zones
<b>Health checks</b>	Automatically checks health of instances and takes them in or out of service
<b>Session stickiness</b>	Route requests to the same instance
<b>Secure sockets layer</b>	Supports SSL offload from web and application servers with flexible cipher support
<b>Monitoring</b>	Publishes metrics to CloudWatch

Scaling this horizontally  
and vertically will get  
us pretty far  
( 10s-100s of thousands )

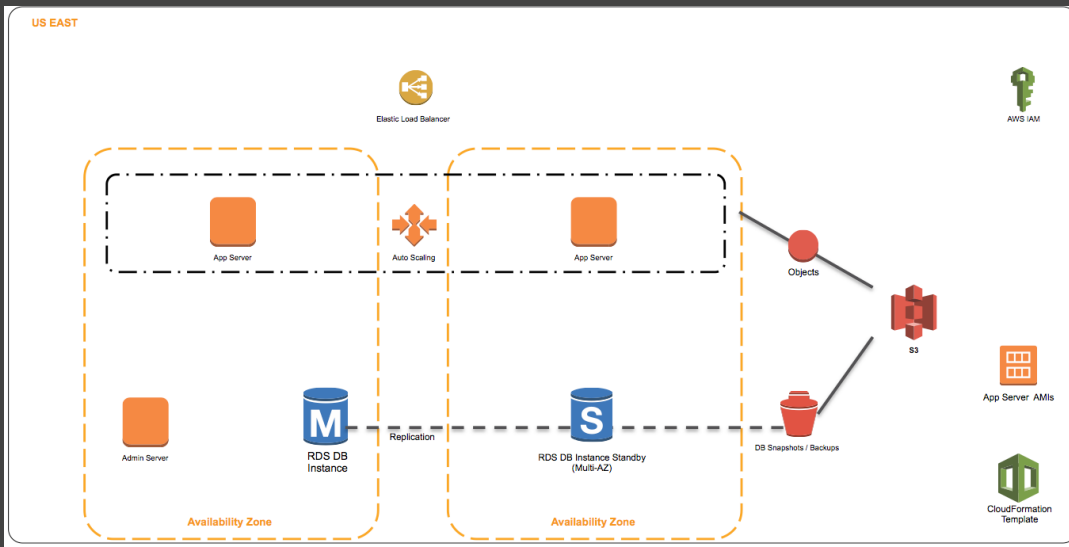


# User > 10ks-100ks:



# Lab 4: HA

- RDS behavior during failover
- Self-healing limitations



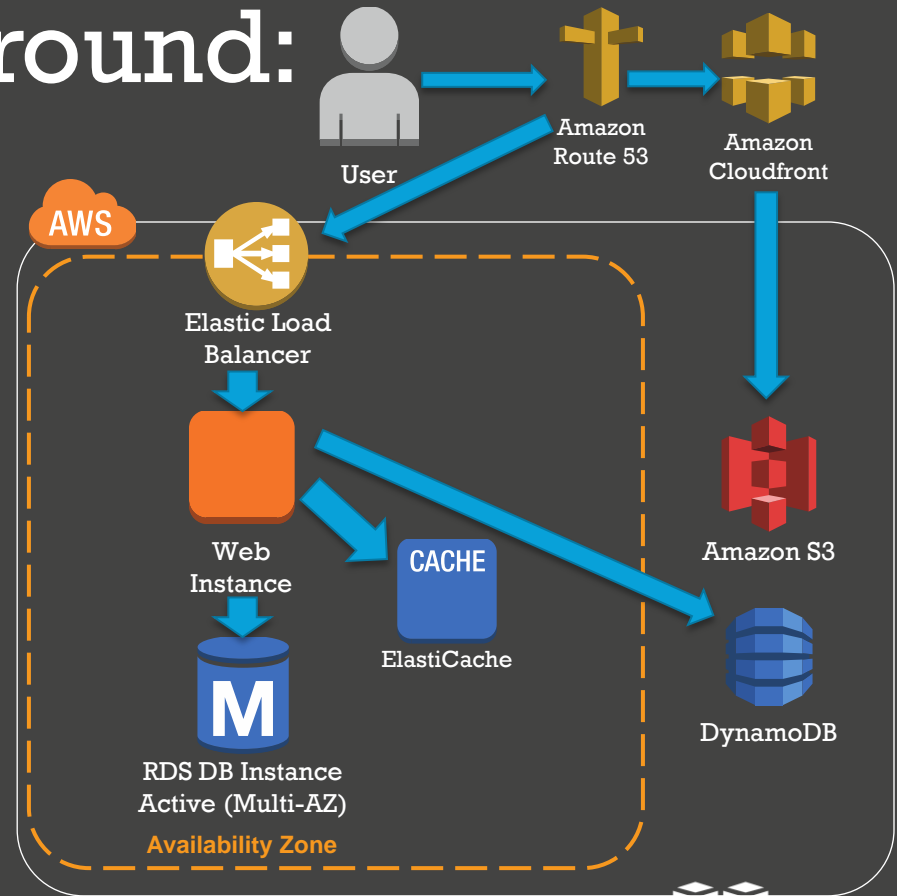
This will take us pretty far  
honestly, but we care about  
performance and efficiency,  
so let's clean this up a bit



# Shift some load around:

Let's lighten the load on our web and database instances:

- Move static content from the Web Instance to S3 and CloudFront
- Move session/state and DB caching to ElastiCache or DynamoDB



# Working with S3 – Amazon Simple Storage Service

- Object based storage for the web
- 11 9s of durability
- Good for things like:
  - Static assets ( css, js, images, videos )
  - Backups
  - Logs
  - Ingest of files for processing
- “Infinitely scalable”
- Supports fine grained permission control
- Ties in well with CloudFront
- Ties in with EMR
- Acts as a logging endpoint for S3/CloudFront/Billing
- Supports Encryption at transit and at rest
- Reduced Redundancy 1/3 cheaper
- Glacier for super long term storage



**Jeff Barr** @jeffbarr

18 Apr

Announced at AWS Summit - Amazon S3 now holds 2 trillion objects, processes 1.1 million requests / second: [bit.ly/ZBN5k2](https://bit.ly/ZBN5k2)  
[#awssummit](#)



# DynamoDB

- Provisioned throughput NoSQL database
- Fast, predictable performance
- Fully distributed, fault tolerant architecture
- Considerations for non-uniform data



Feature	Details
<b>Provisioned throughput</b>	Dial up or down provisioned read/write capacity.
<b>Predictable performance</b>	Average single digit millisecond latencies from SSD-backed infrastructure.
<b>Strong consistency</b>	Be sure you are reading the most up to date values.
<b>Fault tolerant</b>	Data replicated across Availability Zones.
<b>Monitoring</b>	Integrated to CloudWatch.
<b>Secure</b>	Integrates with AWS Identity and Access Management (IAM).
<b>Elastic MapReduce</b>	Integrates with Elastic MapReduce for complex analytics on large datasets.



# ElastiCache

- Hosted Memcached
  - Speaks same API as traditional open source memcached
- Scale from one to many nodes
- Self healing ( replaces dead instance )
- Very fast ( single digit ms speeds usually (or less) )
- Local to a single AZ
  - So need to run different clusters across different AZs
- Data is only in memory, so not persistent
- Use AWS's Auto Discovery client to simplify clusters growing and shrinking without affecting your application

CACHE



Now that our Web tier is  
much more lightweight, we  
can revisit the beginning of  
our talk...

# Auto-Scaling!



# Auto-Scaling




Trigger auto-scaling  
policy

Automatic resizing of compute  
clusters based on demand




Amazon  
CloudWatch



Feature	Details
Control	Define minimum and maximum instance pool sizes and when scaling and cool down occurs.
Integrated to Amazon CloudWatch	Use metrics gathered by CloudWatch to drive scaling.
Instance types	Run Auto Scaling for On-Demand and Spot Instances. Compatible with VPC.

```
as-create-auto-scaling-group MyGroup  
--launch-configuration MyConfig  
--availability-zones us-east-1a  
--min-size 4  
--max-size 200
```

# Typical weekly traffic to Amazon.com



Sunday

Monday

Tuesday

Wednesday

Thursday


Friday

Saturday



# Typical weekly traffic to Amazon.com

Provisioned capacity



Sunday

Monday

Tuesday

Wednesday


Thursday

Friday

Saturday



# November traffic to Amazon.com



November

# November traffic to Amazon.com

Provisioned capacity




November



# November traffic to Amazon.com

Provisioned capacity


76%



November

24%

# November traffic to Amazon.com



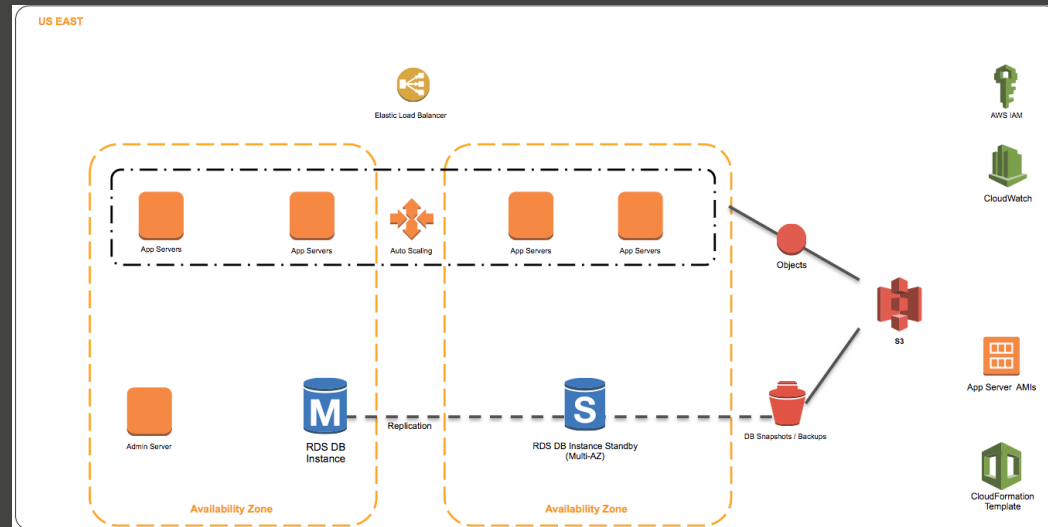
November

Auto-Scaling lets you do this!



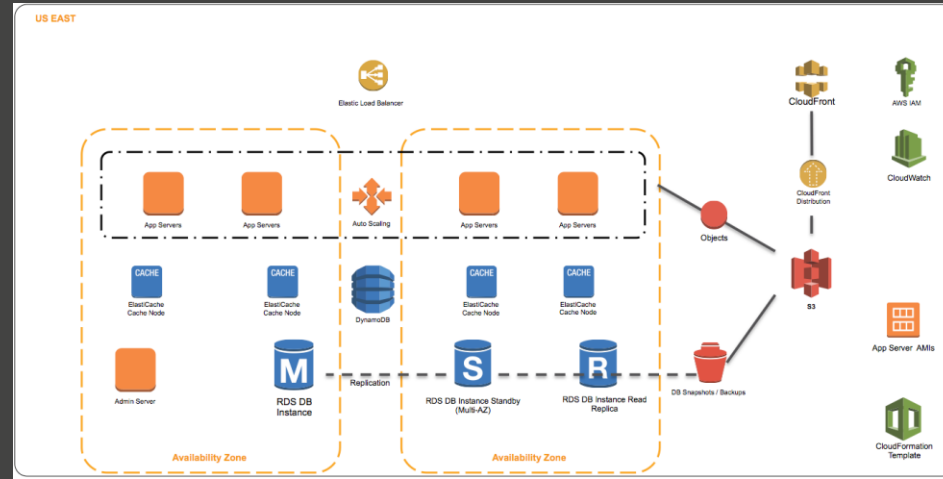
# Lab 5 Review: Auto Scaling

- Approach to Auto Scaling Policies
- Approaches to Load Simulation



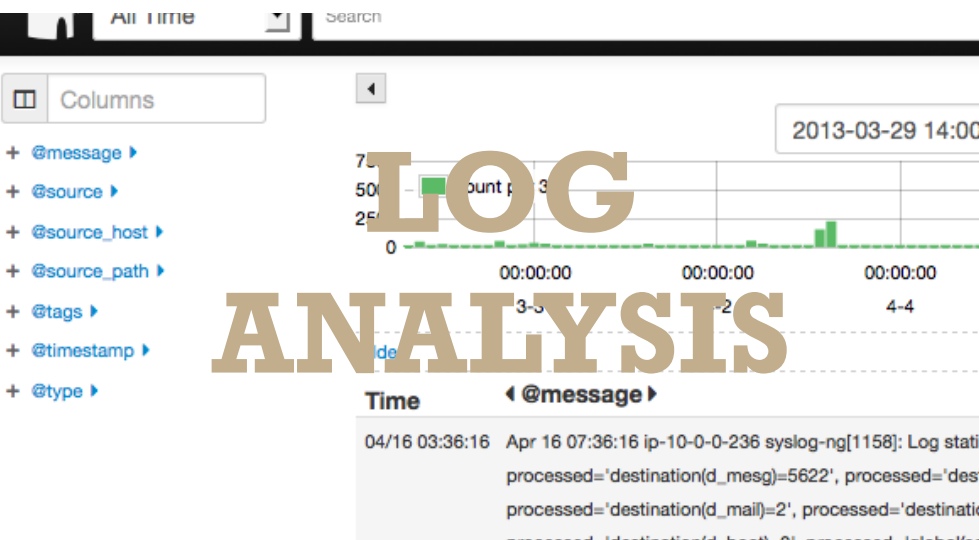
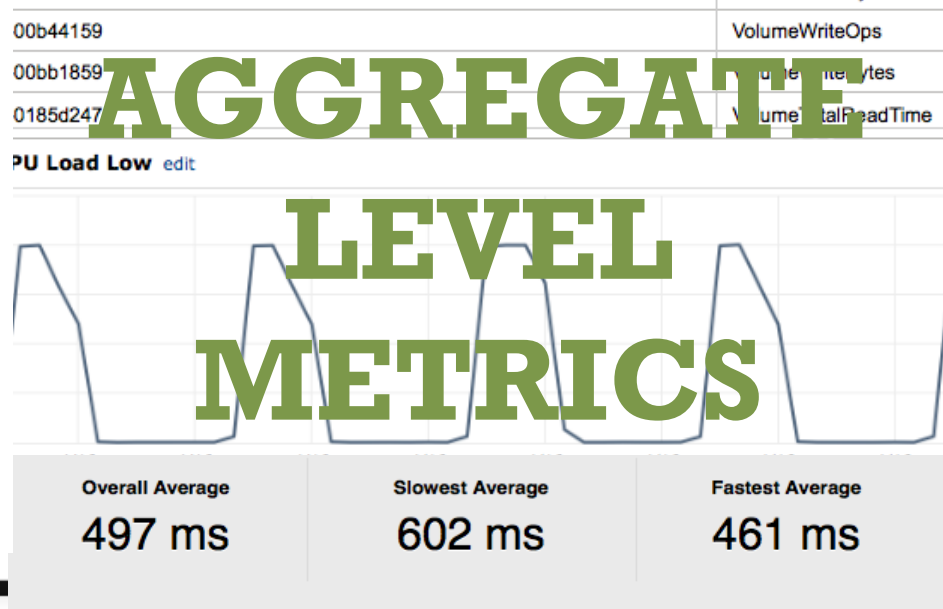
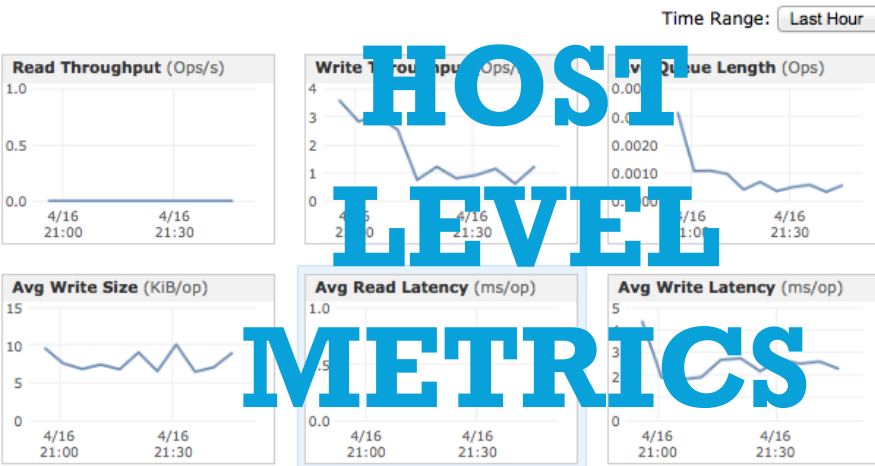
# Option Lab 7: Caching

- static content CloudFront consideration
- MemCached and Elastic Cache
- DynamoDB implementation



Not having proper  
monitoring/metrics is like  
flying a plane with an eye  
mask on in a thunderstorm.  
Oh and your wing is on fire.






# AWS Marketplace & Partners Can Help

- Customer can find, research, buy software
- Simple pricing, aligns with EC2 usage model
- Launch in minutes
- Marketplace billing integrated into your AWS account
- 700+ products across 20+ categories

boundary



DATADOG



Learn more at: [aws.amazon.com/marketplace](https://aws.amazon.com/marketplace)





# Option Lab 8: Monitoring

- Custom CloudWatch Metrics
- Third party Monitoring

# Next steps?

READ! –

- [aws.amazon.com/documentation](https://aws.amazon.com/documentation)
- [aws.amazon.com/architecture](https://aws.amazon.com/architecture)
- [aws.amazon.com/start-ups](https://aws.amazon.com/start-ups)



# Next steps?

START USING AWS –

[aws.amazon.com/free/](https://aws.amazon.com/free/)



# Next steps?

## ASK FOR HELP!

- [forums.aws.amazon.com](https://forums.aws.amazon.com)
- [aws.amazon.com/support](https://aws.amazon.com/support)
- Your local account manager

# THANKS FOR LISTENING!

- [jman@amazon.com](mailto:jman@amazon.com)

